

Digital IC Design

Lecture 7:

Energy-Efficient Architecture

黃柏蒼 Po-Tsang (Bug) Huang

bughuang@nycu.edu.tw


International College of Semiconductor Technology
National Chiao Tung Yang Ming University

國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Energy–Delay Optimization

- Maximize throughput for given energy or Minimize energy for given throughput



Other important metrics: Area, Reliability, Reusability

Design Abstraction Stack Layers

- A very rich set of design parameters to consider!
It helps to consider options in relation to their abstraction layer.

System/Application	Choice of algorithm
Software	Concurrency, Power Control
(Micro-)Architecture	Parallel, Pipeline, Configurable
Logic	logic family, standard cell, custom
Circuit	Sizing, supply, thresholds
Device	Bulk, SOI

Circuit Optimization Framework


Minimize Energy (V_{DD} , V_{TH} , W)
subject to Delay (V_{DD} , V_{TH} , W) $\leq D_{con}$

Constraints

$$V_{DD_min} < V_{DD} < V_{DD_max}$$

$$V_{TH_min} < V_{TH} < V_{TH_max}$$

$$W_{_min} < W$$




■ Reference case

◆ D_{min} Sizing @ (V_{DD_max} , V_{TH_ref})

Impact of Variations

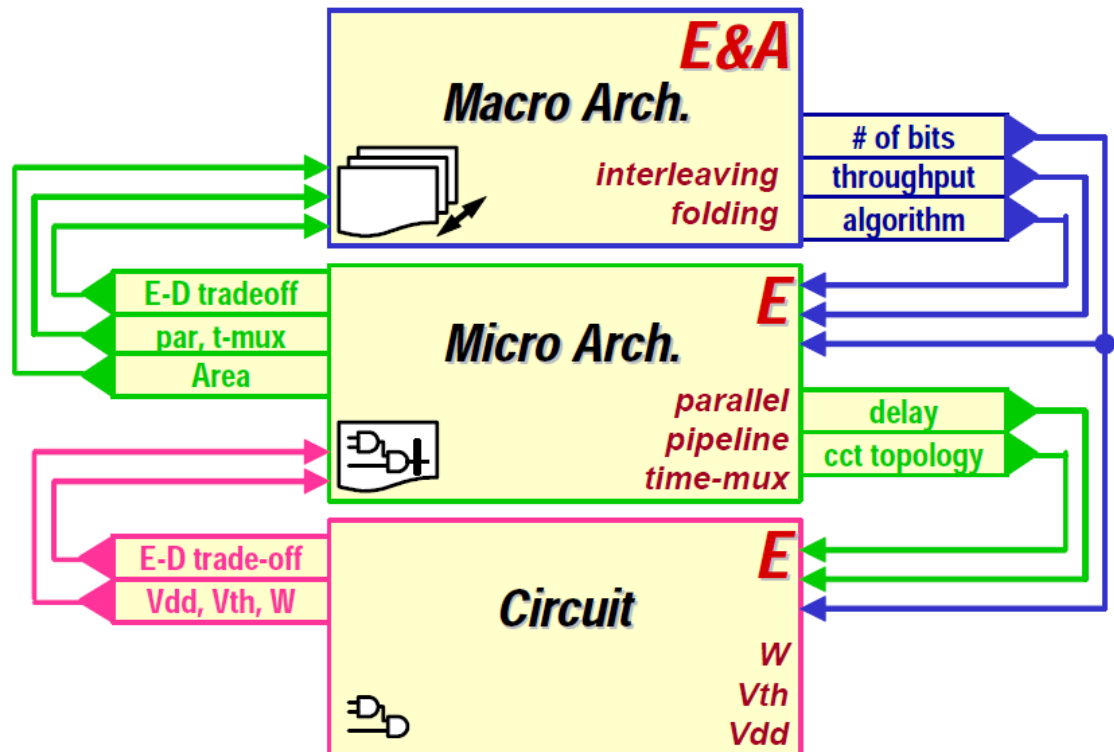
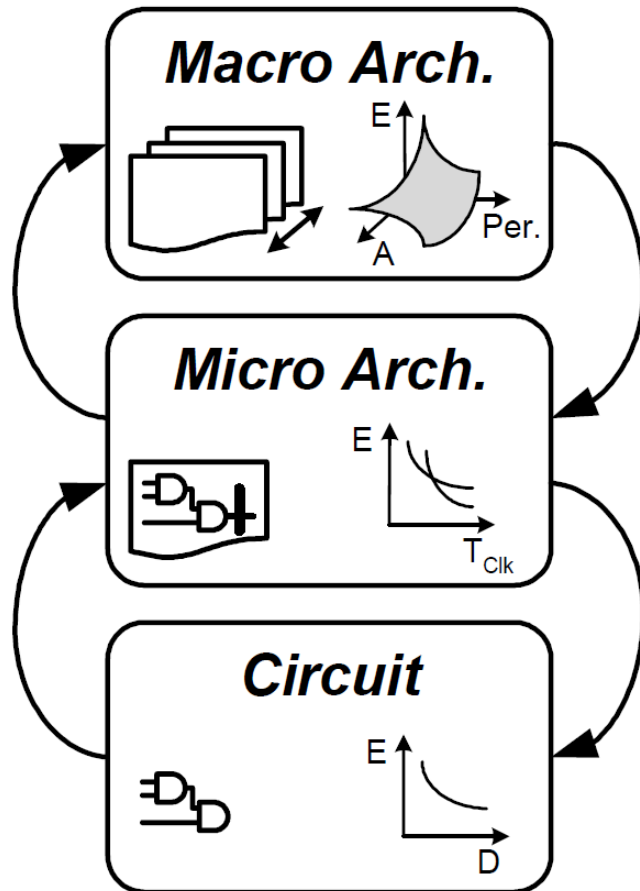
- Downsizing and/or lowering the supply on the critical path lowers the operating frequency
- Downsizing non-critical paths reduces energy for free, but
 - ◆ Narrows down the path–delay distribution
 - ◆ Increases impact of variations, impacts robustness



Optimizing Power at Architecture/System

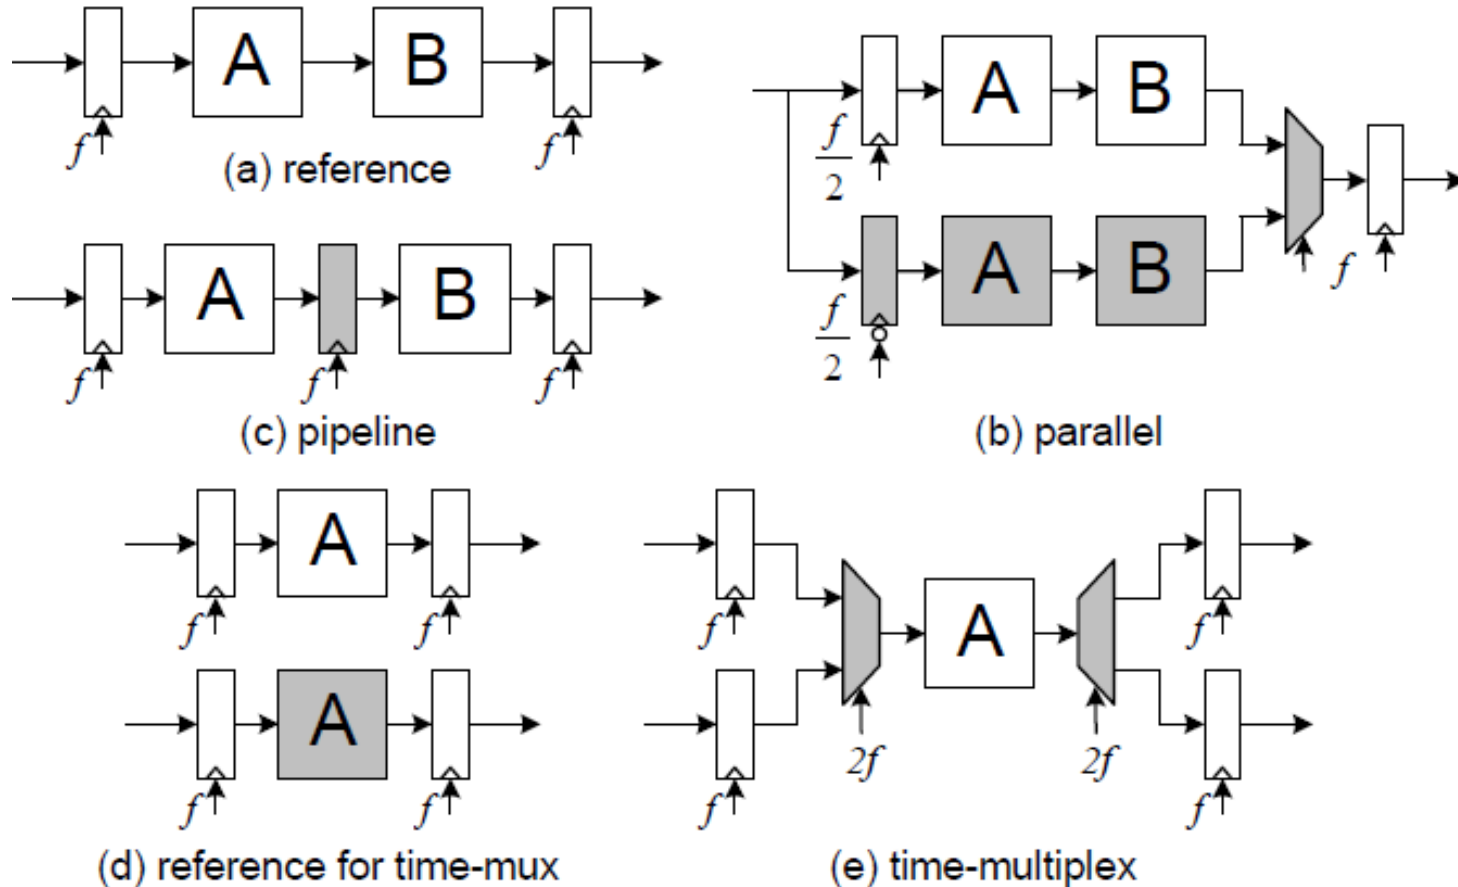
- Optimizations at the architecture or system level can enable more effective power minimization at the circuit level (while maintaining performance), such as
 - ◆ Enabling a reduction in supply voltage
 - ◆ Reducing the effective switching capacitance for a given function (physical capacitance, activity)
 - ◆ Reducing the switching rates
 - ◆ Reducing leakage
- Optimizations at higher abstraction levels tend to have greater potential impact
 - ◆ While circuit techniques may yield improvements in the 10–50% range, architecture and algorithm optimizations have reported power reduction by orders of magnitude

Cross-Layer Optimization



Basic Micro-Architecture Techniques

■ Parallelism, pipeline, time-multiplexing





Reference Datapath

■ Critical-path delay

- ◆ Total capacitance being switching = C_{ref}
- ◆ $VDD = V_{DD_ref}$
- ◆ Power for reference datapath

$$P_{ref} = f_{ref} \cdot C_{ref} \cdot V_{DD_ref}^2$$



Area = $636 \times 833 \mu^2$

Concurrency: trading off frequency for area to reduce power

Parallel Datapath Architecture

- The clock rate of a parallel datapath can be reduced by half with the same throughput



- ◆ $f_{\text{par}} = f_{\text{ref}}/2$

- ◆ $V_{\text{DD_par}} = V_{\text{DD_ref}}/1.7$

- ◆ $C_{\text{par}} = 2.15 * C_{\text{ref}}$




$$P_{\text{par}} \approx 0.36 P_{\text{ref}}$$



Area = 1476 x 1219 μ^2

A Parallel Implementation

- Running slower reduces required supply voltage yields quadratic reduction in power




$$\begin{aligned} f_{\text{par}} &= f_{\text{ref}}/2 \\ C_{\text{par}} &= (2 + ov_{\text{par}}) \cdot C_{\text{ref}} \\ V_{\text{DD par}} &= \epsilon_{\text{par}} \cdot V_{\text{DD ref}} \end{aligned}$$

Almost cancels

$$P_{\text{par}} = \epsilon_{\text{par}}^2 \cdot \left(\frac{2 + ov_{\text{par}}}{2} \right) \cdot P_{\text{ref}}$$

Example for Parallelism in 90nm






Assuming
 $ov_{\text{par}} = 7.5\%$


$$P_{\text{par}} = 0.66^2 \cdot \frac{2.15}{2} \cdot P_{\text{ref}} = 0.47 P_{\text{ref}}$$

$$P_{\text{par } 4} = 0.52^2 \cdot \frac{4.3}{4} \cdot P_{\text{ref}} = 0.29 P_{\text{ref}}$$

Parallelism Adds Latency

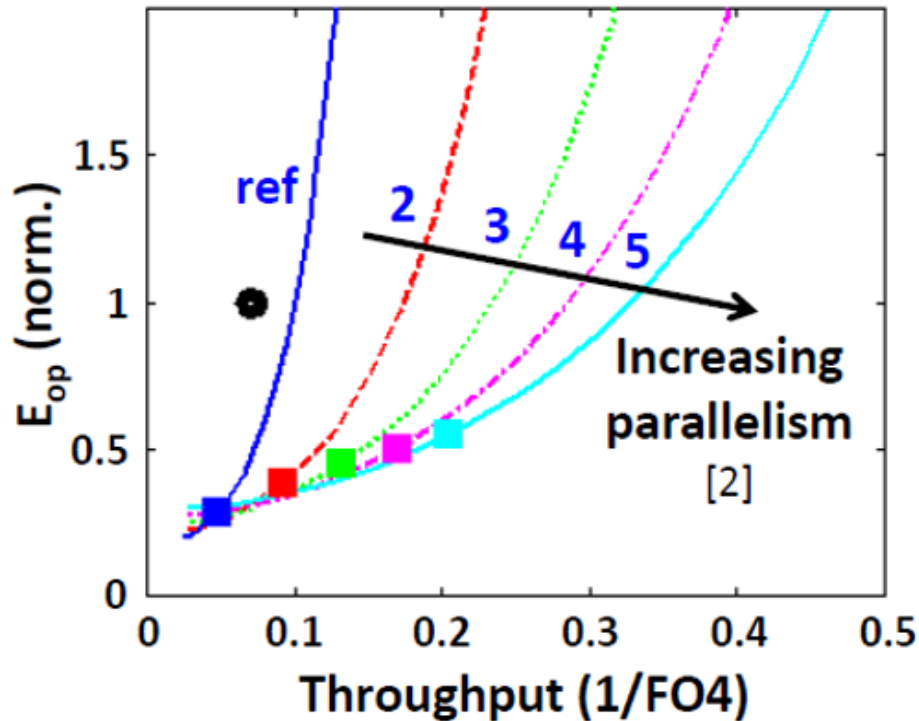


Level of parallelism $P = 2$



Increasing Level of Parallelism

- Improve throughput for the same energy
- Improve energy for the same throughput




$$\text{Area: } A_{\text{par}} \approx N * A_{\text{ref}}$$

Cost: increased area

More Parallelism is NOT Always Better

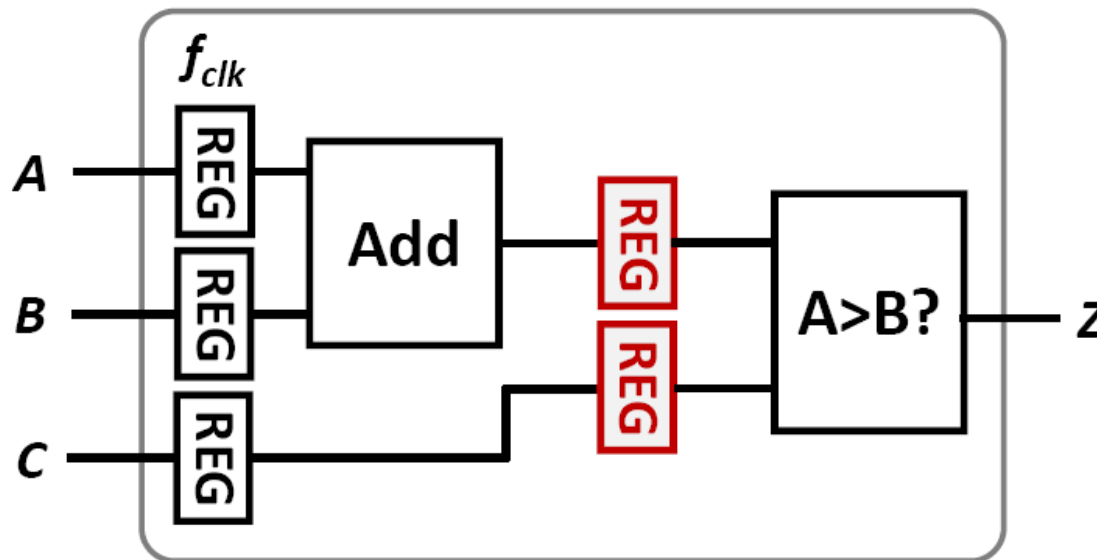
- Leakage and overhead start to dominate at high levels of parallelism, causing minimum energy (dot) to increase
- Optimum voltage also increases with parallelism



$$E_{\text{total}} = E_{\text{SW}} + N \times E_{\text{leak}} + E_{\text{overhead}}$$


Pipelined Datapath Architecture

- Critical-path delay $\rightarrow \max(t_{adder}, t_{comparator})$
 - ◆ Keeping clock rate constant: $f_{pipe} = f_{ref}$
 - ◆ Voltage can be dropped $\rightarrow V_{DD_pipe} = V_{DD_ref} / 1.7$
 - ◆ Capacitance slightly higher: $C_{pipe} = 1.15 \cdot C_{ref}$
 - ◆ $P_{pipe} = f_{ref} \cdot (1.15 \cdot C_{ref}) \cdot (V_{DD_ref} / 1.7)^2 \approx 0.39 \cdot P_{ref}$



A Pipelined Implementation

- Shallower logic reduces required supply voltage



$$\begin{aligned} f_{pipe} &= f_{ref} \\ C_{pipe} &= (1 + ov_{pipe}) \cdot C_{ref} \\ V_{DDpipe} &= \varepsilon_{pipe} \cdot V_{DDref} \end{aligned}$$

$$P_{pipe} = \varepsilon_{pipe}^2 \cdot (1 + ov_{pipe}) \cdot P_{ref}$$

Assuming

$ov_{pipe} = 10\%$

$$P_{pipe} = 0.66^2 \cdot 1.1 \cdot P_{ref} = 0.48 P_{ref}$$

$$P_{pipe4} = 0.52^2 \cdot 1.1 \cdot P_{ref} = 0.29 P_{ref}$$

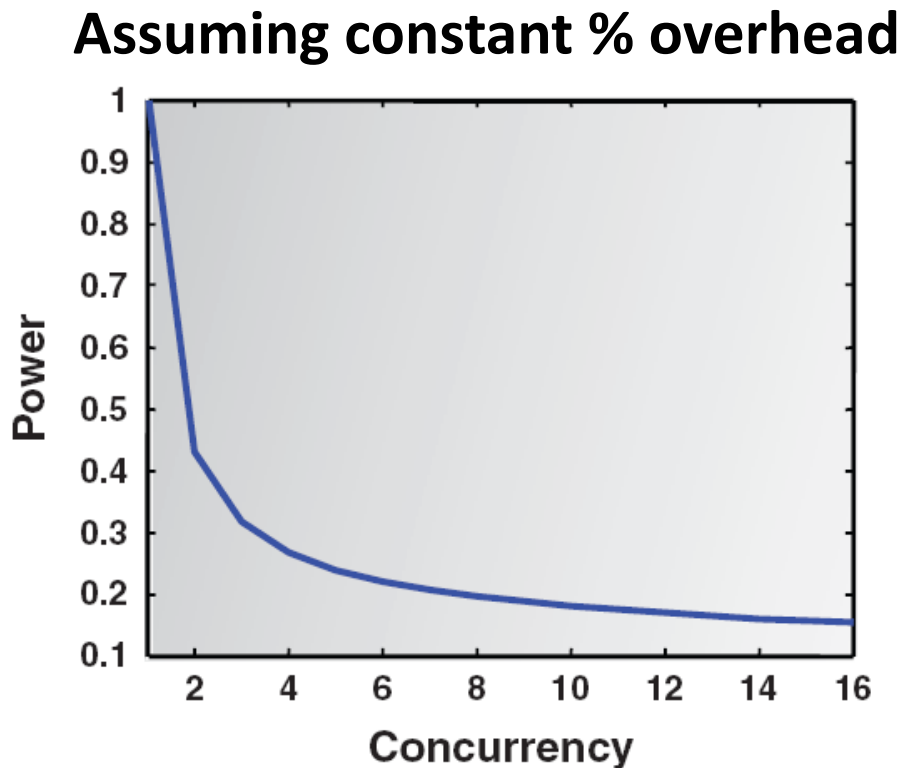
Architecture Summary (Simple Datapath)

- Pipelining and parallelism relax performance of a datapath, which allows voltage reduction and results in power savings
- Pipelining has less overhead than parallelism, but is generally harder to implement (involves finding convenient logic cut-sets)

Architecture type	Voltage	Area	Power
Reference datapath (no pipelining or parallelism)	5 V	1	1
Pipelined datapath	2.9 V	1.3	0.39
Parallel datapath	2.9 V	3.4	0.36
Pipeline-Parallel	2.0 V	3.7	0.2


Increasing Use of Concurrency

- Can combine parallelism and pipelining to drive *supply voltage* down
- But close to process threshold, overhead of excessive concurrency starts to dominate

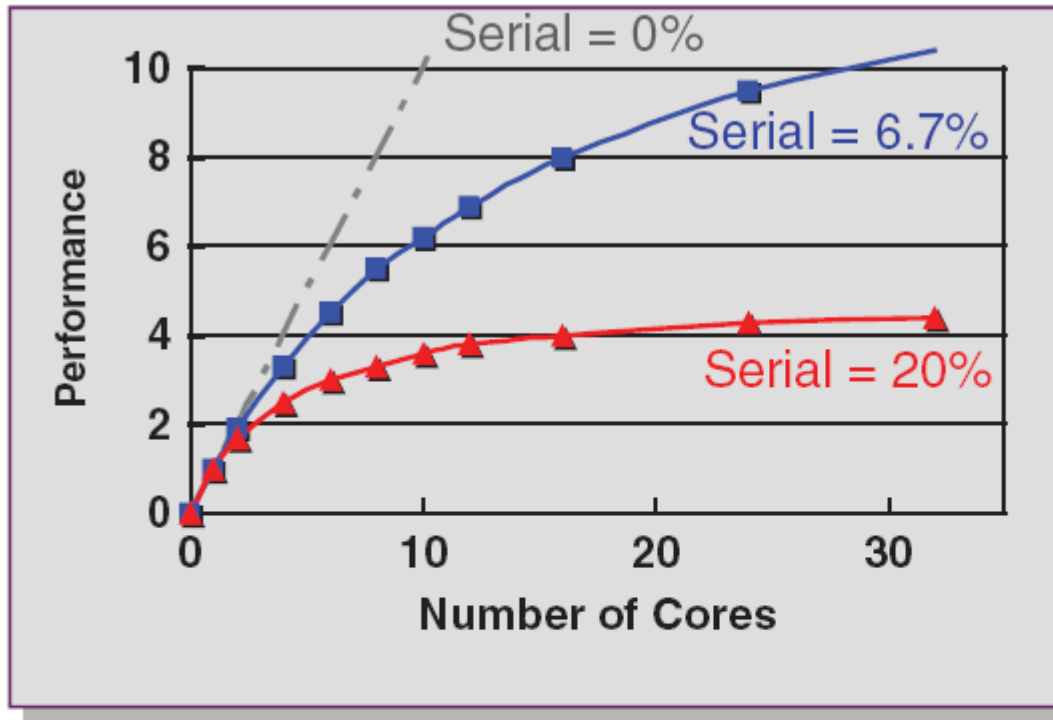


Increasing Use of Concurrency

- Must consider leakage



The Quest for Concurrency




Amdahl's Law:

$$Speed-up = \frac{1}{Serial + \frac{1-Serial}{N}}$$


Performance/Area in Multi-Core

13 mm, 100 W, 48 MB Cache, 4B Transistors, in 22 nm




12 Cores



48 Cores


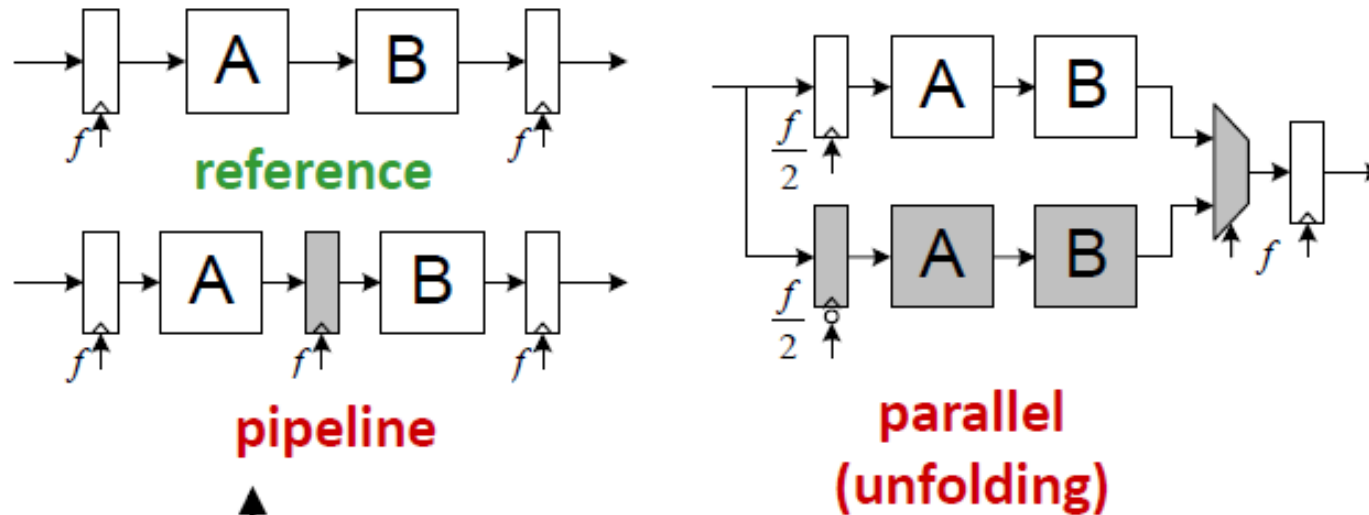


144 Cores





[Courtesy : S. Borkar, Intel, 2006]

Parallelism and Pipelining




Time Multiplexing




Data-Stream Interleaving

PE = recursive operation (feedback)




Large area

Interleaving Architecture




- Reduced area
- P/S overhead
- Pipelined


PE Performs Recursive Operation



➤ Interleave = up-sample & pipeline

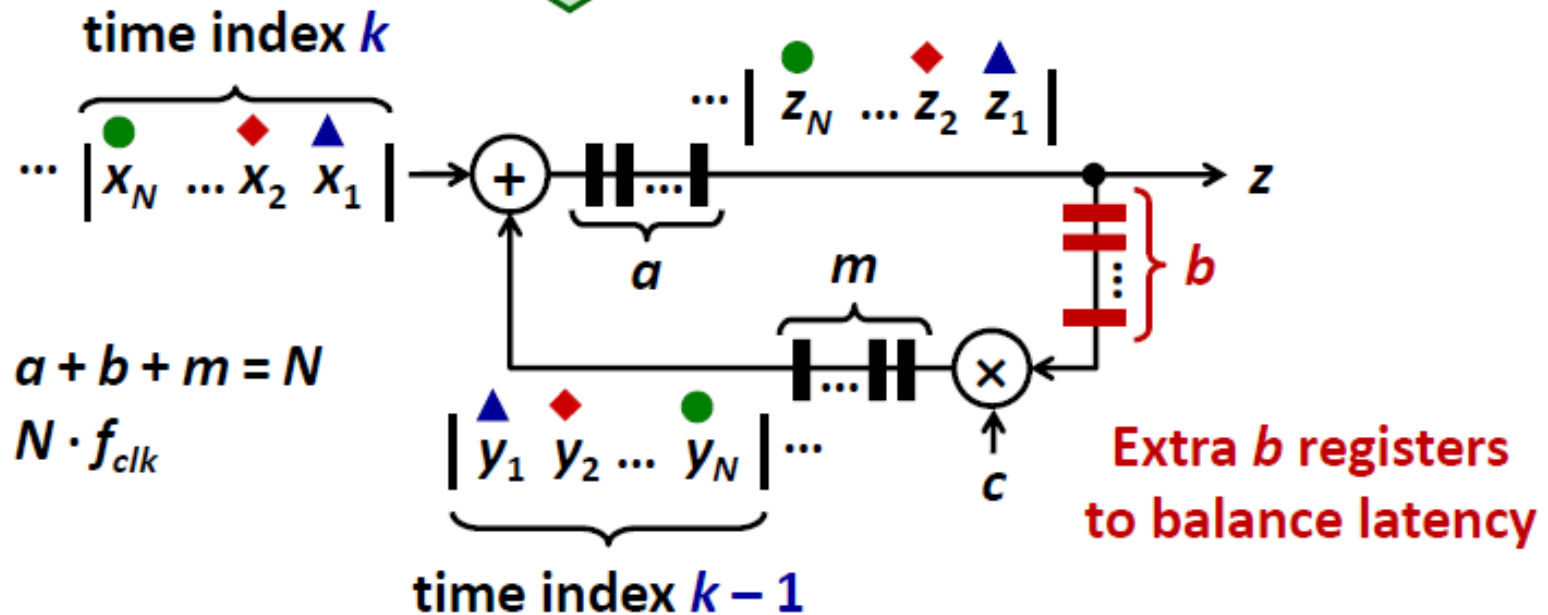


Data-Stream Interleaving Example




Recursive operation:
 $z(k) = x(k) + c \cdot z(k-1)$

\downarrow N data streams: x_1, x_2, \dots, x_N




Folding

PE = recursive operation



- PEs in serial
- Large area

Folded Architecture




- Reduced area
- P/S overhead
- Pipelined

Folding Example


■ Folding = up-sampling & pipelining

- ◆ Reduced area (shared datapath logic)




Area Benefit of Interleaving and Folding

- Area: $A = A_{logic} + A_{registers}$
- Interleaving or folding of level N
 - ◆ $A = A_{logic} + N \times A_{registers}$
- Timing and Energy stay the same



Architecture Transformations


- Procedure: move toward desired E-D point while minimizing area



Parallelism & Pipelining Transformations

■ Parallelism & Pipelining


◆ reduce Energy, increase Area



Time Multiplexing Transformations

■ Time multiplexing


◆ increase Energy, reduce Area



Interleaving & Folding Transformations

■ Interleaving & Folding (time-mux + pipeline)

◆ \approx const Energy, reduce Area



Some Energy-Inspired Design Guidelines

■ **For maximum performance**

- ◆ Maximize use of concurrency at the cost of area

■ **For given performance**

- ◆ Optimal amount of concurrency for minimum energy


■ **For given energy**

- ◆ Least amount of concurrency that meets performance goals

■ **For minimum energy**

- ◆ Solution with minimum overhead (that is – direct mapping between function and architecture)

Improving Computational Efficiency



Implementations for a given function maybe inefficient and can often be replaced with more efficient versions **without penalty in energy or delay**

■ Inefficiencies arise from:

- ◆ Over-dimensioning or over-design
- ◆ Generality of function
- ◆ Design methodologies
- ◆ Limited design time
- ◆ Need for flexibility, re-use, and programmability